

C introduction

# Dynamic memory management

# Contents

# Runtime conditions

Static C arrays are great if you already know at *compile time* how many elements you will need later at *runtime*.

Seriously, how often is this the case?

# Runtime conditions

Static C arrays are great if you already know at *compile time* how many elements you will need later at *runtime*.

Seriously, how often is this the case?

*Never.* Unless your program does not take *any* user input, you cannot determine how much data you will have to store.

This is where dynamic memory management comes into play.

# Explicit allocation


All the variables and arrays you have used so far were placed in memory automatically. In dynamic memory management, you have to *allocate* parts memory to identifiers on your own.

There are four functions in the standard library that do almost all the work for you:

- ▶ *malloc()*: Allocate a block of memory
- ▶ *calloc()*: Allocate a block of memory and initialize it
- ▶ *realloc()*: Alter the size of a block of memory
- ▶ *free()*: Release a block of memory


They are declared in *stdlib.h*.

## A closer look at memory




All local variables of functions are placed at the *stack*.  
It grows and shrinks as variables are declared and functions return.

## A closer look at memory



Dynamical memory is allocated on the *heap*.  
The example shows a function with two local `int` variables.

## A closer look at memory




```
malloc( sizeof( int ) );
```

Reserves exactly the amount of memory an *int* variable takes.




## A closer look at memory



```
int *new_block = malloc(sizeof(int));
```

The address of that memory block is stored in an *int* pointer.

## A closer look at memory



*malloc()* just needs to know the size of the block it reserves.  
Let us allocate a *long* variable as well.

## *malloc()* in detail

The function declaration might be a little bit confusing:

```
void *malloc(size_t size);
```

- ▶ *size\_t* is an **unsigned integer** type.  
Any positive integer number (e.g. an *int* > 0) will do the job.
- ▶ *size* is the size of the reserved block in **bytes**.  
If you want to use that block *seriously*, pass the size of an actual type (e.g. *sizeof(int)*).
- ▶ A *void* pointer is returned since *malloc()* does not know how you want to use the reserved block. By assigning it to a regular pointer variable it is automatically converted to that type.

## Casting or not casting. . .

Some people claim you had to explicitly *cast* the return value of *malloc()*.

## Casting or not casting. . .

Some people claim you had to explicitly *cast* the return value of *malloc()*.

**This is outdated.**

## Casting or not casting...

Some people claim you had to explicitly *cast* the return value of *malloc()*.

**This is outdated.**

```
int *block = malloc(sizeof(int));
```

has the same result as

```
int *block = (int*) malloc(sizeof(int));
```

while the second one contains a redundant *cast* and if you want to change the type of *block* later, you will have to hit more keys. Consider:

```
int *block = malloc(sizeof *block); /* gold standard */
```

... that is the question

Confused?

... that is the question

## Confused?

Do not typecast the result of `malloc()` & Co.



## Tidying up

Unlike normally declared variables, dynamically allocated storage is not automatically released when the function returns.

```
1 void foo() {  
2     int *bar = malloc(sizeof(int));  
3 }
```

With the pointer *bar* being removed from the stack, we have no reference on its allocated memory and those four bytes are blocked forever!

```
free(void *ptr);
```

Pass any pointer to previously allocated memory to *free()* and it gets released. If you pass pointers on other things, undefined behaviour occurs (most likely program crashes).

## Reserving large chunks

To get a dynamic array of a certain *type* and *length*, you have to

- ▶ Pass the block size  $length * sizeof(type)$  to `malloc()`
- ▶ Assign the return value to a pointer to *type*

*int* array with 42 elements:

```
int *field = malloc(42 * sizeof(*field));
```

Since the size of your dynamically allocated array is unknown at compile time, you cannot use `sizeof` to get its length. Save it in its own variable!

With the help of pointer arithmetic, you can use the dynamic array like a "normal" one.

## The fancy alternative

```
void *calloc(size_t nmemb, size_t size);
```

- ▶ Allocates a block of  $nmemb * size$  bytes, where  $nmemb$  is supposed to be the array's length and  $size$  the size of its type.
- ▶ The whole block is filled with 0s

```
int field_length = 42;  
int *field = malloc(field_length * sizeof(*field));  
for (int i = 0; i < field_length; i++)  
    field[i] = 0;
```

↓ Feel the difference ↓

```
int field_length = 42;  
int *field = calloc(field_length, sizeof(*field));
```

# Resizing arrays

Now we come to the point that motivated us to use dynamic arrays:

```
void *realloc(void *ptr, size_t size);
```

- ▶ *ptr* is a pointer to a dynamically allocated memory block
- ▶ *size* is the wanted new size of the memory block
- ▶ The return value is a pointer to the resized block

Note that the new *size* can be greater or smaller than the old one!

- ▶ If it's smaller, you may lose some data at the end of the block
- ▶ If it's greater, the block may be at a different location in the memory  
→ *ptr* is freed then, also the additional bytes are not initialized

## Clean up your code

Passing arrays between functions can be complicated if you store the pointer and the length separately.

Do you remember a way to keep different things together?

## Clean up your code

Passing arrays between functions can be complicated if you store the pointer and the length separately.

Do you remember a way to keep different things together?

```
struct int_array {  
    int *field;  
    int length;  
}
```

This allows you to use the *struct int\_array* as a single argument or return value. Even better: pass a pointer on that structure.

## Strings from pointers to *char*

By handling strings as dynamic *char* arrays you can alter their size which is needed for many operations on them.

- ▶ *strlen()* returns the actual length of a string (up to '`\0`' character)
- ▶ *strncpy()* copies a string into a dynamically allocated block

```
char conststr[] = "Hello";           /* not of much use */
int bufsize = strlen(conststr) + 1;  /* add '\0' char */
char *str = calloc(bufsize, sizeof(*str));
str = strncpy(str, conststr, bufsize); /* ready to go */
```

These functions and others are declared in *string.h*.

```
$ man string.h
```

# Exercises

- ▶ You are now able to solve tasks 26, 27 and 28.
- ▶ You additionally are now able to visit the advanced C course.